

Settlement Algebra and Verified Constant-Product Invariants for Batch DEX Execution

(draft; authors TBD)

January 30, 2026

Abstract

We present a small, compositional mathematical model for decentralized exchange (DEX) execution and prove key safety invariants for constant-product market makers (CPMMs), including a closed-form expression for the *exact* per-swap increase in the product invariant. The development is mechanized in Lean 4/Mathlib.

The paper has three main layers: (i) an additive *settlement algebra* capturing token flows as a commutative group, with a homomorphism Δ to a scalar “net flow”; (ii) a correctness-by-construction objective for batch auctions maximizing executed volume and surplus in lexicographic order; and (iii) a CPMM state-transition model with verified K -monotonicity and an exact “ K -gap” remainder theorem, lifted compositionally to batches and bridged to the settlement/objective layers by a refinement lemma.

1 Overview and scope

This draft explains the mathematics and formal logic embodied by the Lean files:

- `lean-mathlib/Proofs/SettlementAlgebra.lean`: additive settlement model and the homomorphism Δ .
- `lean-mathlib/Proofs/BatchOptimality.lean`: batch objective (A, B) (volume, surplus) and lexicographic comparison.
- `lean-mathlib/Proofs/CPMMInvariants.lean`: CPMM safety theorems, including the closed-form K -gap.
- `lean-mathlib/Proofs/CPMMSettlement.lean`: bridge from CPMM swaps to settlements.
- `lean-mathlib/Proofs/BatchCPMMUnification.lean`: unification: batch settlement safety, batch optimality additivity, and compositional K theorems for sequential swap batches.

Important semantic caveat. Throughout, we distinguish:

- *Scalar net-flow* $\Delta(s) = \Delta x + \Delta y$ which adds different token units and is therefore *not* a notion of economic value conservation; and
- the *CPMM invariant* $K = x \cdot y$ which is the standard, economically meaningful conservation law for constant-product pools.

The Lean development makes this distinction explicit and proves K -monotonicity as the primary invariant.

2 Preliminaries

We write $\mathbb{N} = \{0, 1, 2, \dots\}$ and \mathbb{Z} for the integers. For $d \in \mathbb{N}$ with $d > 0$, Euclidean division states that for every $N \in \mathbb{N}$ there exist unique $q, r \in \mathbb{N}$ such that

$$N = dq + r \quad \text{and} \quad 0 \leq r < d, \tag{1}$$

where $q = \lfloor N/d \rfloor$ and $r = N \bmod d$ [2].

3 Settlement algebra

3.1 Definition

Definition 1 (Settlement). Define the settlement space as

$$\text{Settlement} := \mathbb{Z} \times \mathbb{Z}.$$

An element $s = (\Delta x, \Delta y)$ represents signed changes to two pool reserves. Positive means inflow to the pool, negative means outflow.

Definition 2 (Composition and identity). Define composition as componentwise addition:

$$(\Delta x_1, \Delta y_1) \circ_s (\Delta x_2, \Delta y_2) := (\Delta x_1 + \Delta x_2, \Delta y_1 + \Delta y_2).$$

The identity settlement is $0 := (0, 0)$.

Proposition 1 (Abelian group structure). $(\text{Settlement}, \circ_s, 0)$ is an abelian group (indeed an AddCommGroup in Mathlib).

Proof. Componentwise integer addition is associative and commutative, with identity $(0, 0)$ and inverse $-(\Delta x, \Delta y) = (-\Delta x, -\Delta y)$. \square

3.2 Net-flow homomorphism and balance

Definition 3 (Net scalar flow). Define $\Delta : \text{Settlement} \rightarrow \mathbb{Z}$ by

$$\Delta(\Delta x, \Delta y) := \Delta x + \Delta y.$$

Proposition 2 (Homomorphism). Δ is a group homomorphism:

$$\Delta(s_1 \circ_s s_2) = \Delta(s_1) + \Delta(s_2), \quad \Delta(0) = 0, \quad \Delta(-s) = -\Delta(s).$$

Proof. Immediate from distributivity of addition in \mathbb{Z} :

$$\Delta((\Delta x_1 + \Delta x_2, \Delta y_1 + \Delta y_2)) = (\Delta x_1 + \Delta x_2) + (\Delta y_1 + \Delta y_2) = (\Delta x_1 + \Delta y_1) + (\Delta x_2 + \Delta y_2).$$

\square

Definition 4 (Balanced settlements). Call a settlement balanced if $\Delta(s) = 0$. Equivalently,

$$\text{Balanced} := \ker(\Delta) \subseteq \text{Settlement}.$$

Proposition 3 (Kernel is a subgroup). Balanced is an additive subgroup of Settlement.

Proof. Kernels of homomorphisms are subgroups; explicitly: if $\Delta(s_1) = 0$ and $\Delta(s_2) = 0$, then $\Delta(s_1 \circ_s s_2) = 0$, and if $\Delta(s) = 0$ then $\Delta(-s) = 0$. \square

Remark 1 (“Safety” as a scalar predicate). The Lean development also defines a *scalar safety* predicate

$$\text{isSafe}(s) \iff \Delta(s) \geq 0.$$

This is closed under composition by the homomorphism property and monotonicity of $+$ on \mathbb{Z} , but it should not be conflated with economic safety, since it adds different token units.

3.3 Swap-to-settlement embedding

Definition 5 (Swap settlement). *Given natural amounts $\text{amt}_{\text{in}}, \text{amt}_{\text{out}} \in \mathbb{N}$, define the swap settlement (pool perspective)*

$$\text{swap}(\text{amt}_{\text{in}}, \text{amt}_{\text{out}}) := (\text{amt}_{\text{in}}, -\text{amt}_{\text{out}}) \in \mathbb{Z} \times \mathbb{Z}.$$

Then

$$\Delta(\text{swap}(\text{amt}_{\text{in}}, \text{amt}_{\text{out}})) = \text{amt}_{\text{in}} - \text{amt}_{\text{out}}.$$

In particular, if $\text{amt}_{\text{in}} = \text{amt}_{\text{out}}$ then the settlement is balanced.

4 Batch optimality: maximizing volume and surplus

This section formalizes an objective for batch execution that is *correctness-by-construction*: invalid fills (output below minimum) are unrepresentable.

4.1 Intents and valid outcomes

Definition 6 (Intent). *An intent is a pair $i = (a, m) \in \mathbb{N} \times \mathbb{N}$, where a is the input amount and m is the minimum acceptable output.*

Definition 7 (Valid outcome). *For a fixed intent $i = (a, m)$, a valid outcome is either:*

1. *unfilled*, or
2. *filled*(o) for some $o \in \mathbb{N}$ with the proof obligation $o \geq m$.

This is implemented as a dependent inductive type in Lean (`ValidOutcome i`).

4.2 Objective functions

Definition 8 (Volume and surplus). *For intent $i = (a, m)$ and valid outcome o :*

$$\text{volume}(o) := \begin{cases} 0 & \text{if } \text{unfilled} \\ a & \text{if } \text{filled} \end{cases} \quad \text{surplus}(o) := \begin{cases} 0 & \text{if } \text{unfilled} \\ o - m & \text{if } \text{filled with output } o \geq m. \end{cases}$$

Because filled outcomes carry the proof $o \geq m$, the subtraction $o - m$ is well-defined in \mathbb{N} (no underflow).

Definition 9 ((A, B) pair). Define the objective for a single intent/outcome as

$$\text{AB}(o) := (\text{volume}(o), \text{surplus}(o)) \in \mathbb{N} \times \mathbb{N}.$$

For a batch (a finite list of intent/outcome pairs), define the batch score as componentwise sum:

$$\text{AB}_{\text{batch}} := \sum_j \text{AB}(o_j).$$

In Lean this is `batchAB`, computed by a left fold with `pairAdd`.

4.3 Lexicographic comparison

Definition 10 (Lexicographic order). For $p = (p_1, p_2), q = (q_1, q_2) \in \mathbb{N} \times \mathbb{N}$, define

$$p \preceq_{\text{lex}} q \iff (p_1 < q_1) \vee (p_1 = q_1 \wedge p_2 \leq q_2),$$

and strict order $p \prec_{\text{lex}} q$ analogously with $p_2 < q_2$.

Proposition 4 (Total preorder). \preceq_{lex} is reflexive, transitive, and total on $\mathbb{N} \times \mathbb{N}$, and the induced equivalence yields antisymmetry for \preceq_{lex} on pairs.

Proof. Standard case analysis on the first components; see Lean theorems:

- `lexLe_refl`
- `lexLe_trans`
- `lexLe_total`
- `lexLe_antisymm`

□

5 CPMM model and invariants

We model a constant-product pool with reserves $(x, y) \in \mathbb{N}^2$ as in Uniswap V2 [1].

5.1 Zero-fee swap semantics

Definition 11 (Zero-fee output). Given reserves $(x, y) \in \mathbb{N}^2$ and input $a \in \mathbb{N}$, define the (zero-fee) output as

$$q := \left\lfloor \frac{ya}{x+a} \right\rfloor. \quad (2)$$

The updated reserves are

$$x' := x + a, \quad y' := y - q. \quad (3)$$

Lemma 1 (Output is reserve-bounded). For all $x, y, a \in \mathbb{N}$, the output satisfies $q \leq y$.

Proof. Since $a \leq x + a$, we have $ya \leq y(x + a)$. Dividing by the positive denominator $x + a$ gives $\frac{ya}{x+a} \leq y$, and taking floors yields $q \leq y$. □

5.2 The constant-product invariant and the K -gap

Definition 12 (Constant product). Define $K(x, y) := x \cdot y$.

Theorem 1 (K -monotonicity). Under a zero-fee swap update (2)–(3), the product never decreases:

$$K(x', y') \geq K(x, y).$$

Proof (via the exact K -gap identity). We strengthen the claim by showing the exact difference is a Euclidean remainder. Let $d := x + a$ and $N := ya$. By Euclidean division (1), there exist $q, r \in \mathbb{N}$ with $N = dq + r$ and $0 \leq r < d$. By definition $q = \lfloor N/d \rfloor = \lfloor ya/(x + a) \rfloor$ which matches (2). Then:

$$\begin{aligned} K(x', y') &= (x + a)(y - q) \\ &= dy - dq \\ &= dy - (N - r) \\ &= (x + a)y - (ya) + r \\ &= xy + r. \end{aligned}$$

Thus $K(x', y') - K(x, y) = r \geq 0$. This is mechanized as `k_gap_exact` in `CPMMInvariants.lean`. \square

Definition 13 (K -gap remainder). Define the per-swap remainder (the K -gap) as

$$r(x, y, a) := (ya) \bmod (x + a). \quad (4)$$

Then the exact K -gap identity is

$$(x + a)\left(y - \left\lfloor \frac{ya}{x+a} \right\rfloor\right) = xy + r(x, y, a). \quad (5)$$

Corollary 1 (Characterizations). For all $x, y, a \in \mathbb{N}$:

1. $K(x', y') \geq K(x, y)$ (nonnegativity),
2. $K(x', y') < K(x, y) + (x + a)$ (bounded increase, since $r < x + a$),
3. $K(x', y') = K(x, y)$ iff $(x + a) \mid (ya)$ (zero remainder).

Proof. Each follows from (5) and standard properties of mod. See `k_gap_nonneg`, `k_gap_bounded`, `k_gap_zero_iff`. \square

5.3 Batches: telescoping the exact K -gap

Definition 14 (Sequential execution). Fix an initial state (x_0, y_0) and a list of inputs $[a_1, \dots, a_n]$. Define the sequence of states by iterating (2)–(3): $(x_i, y_i) \mapsto (x_{i+1}, y_{i+1})$ using input a_{i+1} .

Theorem 2 (Exact telescoping sum). For the sequential execution above,

$$K(x_n, y_n) = K(x_0, y_0) + \sum_{i=0}^{n-1} r(x_i, y_i, a_{i+1}), \quad (6)$$

where r is the per-step remainder (4).

Proof. Induction on n . The induction step applies the single-step identity (5) at state (x_i, y_i) and accumulates remainders. This is mechanized as `executeBatchSwaps_K_gap_sum`. \square

Corollary 2 (Batch K -monotonicity). Since each remainder is nonnegative, $K(x_n, y_n) \geq K(x_0, y_0)$ for every batch.

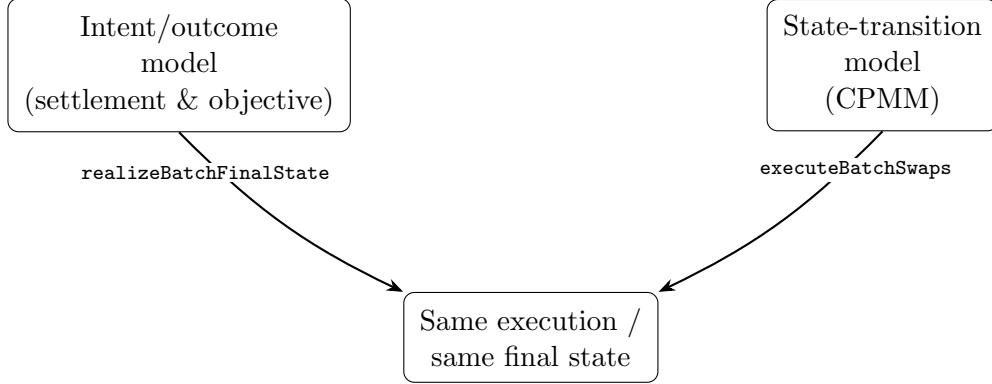


Figure 1: Two views of batch execution that agree by refinement: an intent/outcome view (used for settlement and objective computation) and a reserve state-transition view (used for CPMM invariants).

6 Unification: settlements, objective, and CPMM state

There are two complementary models for batch execution:

1. **Intent/outcome (settlement & objective) model.** A batch is a list of pairs (intent, valid outcome). It induces: (a) an aggregated settlement in `Settlement` (for Δ -analysis) and (b) an aggregated objective score in $\mathbb{N} \times \mathbb{N}$ (for optimization).
2. **State-transition (CPMM) model.** A batch is a list of inputs $[a_1, \dots, a_n]$ executed sequentially on reserves, yielding a final reserve state with verified K properties.

6.1 Aggregating settlements

For each intent/outcome pair, the Lean development defines a per-pair settlement: unfilled contributes 0, and filled contributes $\text{swap}(a, o)$. The batch settlement is the fold (sum) of these settlements. Using the homomorphism property, scalar net-flow composes:

$$\Delta\left(\sum_j s_j\right) = \sum_j \Delta(s_j).$$

Consequently, if each filled execution satisfies $a \geq o$ (a *same-unit* condition), then Δ -safety composes over the batch (`batch_safe_implies_settlement_safe`).

6.2 Refinement bridge and unified safety theorem

The key bridge is that the output-generating function (computing q_i sequentially) reaches the same final state as the pure state-transition fold:

$$\text{realizeBatchFinalState}(s, \vec{a}) = \text{executeBatchSwaps}(s, \vec{a}).$$

This is proved as `realizeBatch_final_state` in `BatchCPMMUnification.lean`.

As a consequence, the CPMM invariants proven for the state-transition model apply to the same execution that feeds settlement/objective computations. The combined statement is captured by `unified_batch_safety`, which simultaneously establishes:

1. $K(\text{final}) \geq K(\text{initial})$,
2. $K(\text{final}) = K(\text{initial}) + \sum \text{remainders as in (6)}$,
3. reserve-in accounting: $x_{\text{final}} = x_{\text{initial}} + \sum_i a_i$.

7 Mechanization notes

The results above are mechanized in Lean 4/Mathlib (no axioms and no `sorry`) [3]. To type-check/build the Lean project:

```
cd lean-mathlib
lake build
```

8 Limitations and future work

This development is deliberately minimal and focuses on arithmetic correctness and composition. Several important directions remain:

- Multi-asset generalization: settlements as \mathbb{Z}^n vectors; richer notions of value conservation.
- Full batch clearing optimality: argmax existence over finite candidate sets and deterministic tie-breaking.
- Multi-hop routing and cross-pool execution: lifting K -style invariants across routing graphs.
- Overflow-aware models and byte-level constraints matching on-chain execution environments.
- Expanded fee models and protocol-specific mechanics (dynamic fees, rebates, TWAP/oracle guards, etc.).

References

- [1] Hayden Adams, Noah Zinsmeister, and Dan Robinson. Uniswap V2. Whitepaper, 2020. Available at <https://uniswap.org/whitepaper.pdf>.
- [2] Euclid. Elements (Book VII): Euclidean division. *Classical source*, 300. Standard statement: for $N, d \in \mathbb{N}$ with $d > 0$, there exist unique q, r s.t. $N = dq + r$ and $0 \leq r < d$.
- [3] Leonardo de Moura and Sebastian Ullrich. *The Lean 4 Theorem Prover and Programming Language*. Lean project, 2023. Reference manual and documentation.